Consulting Craft & the Case Method

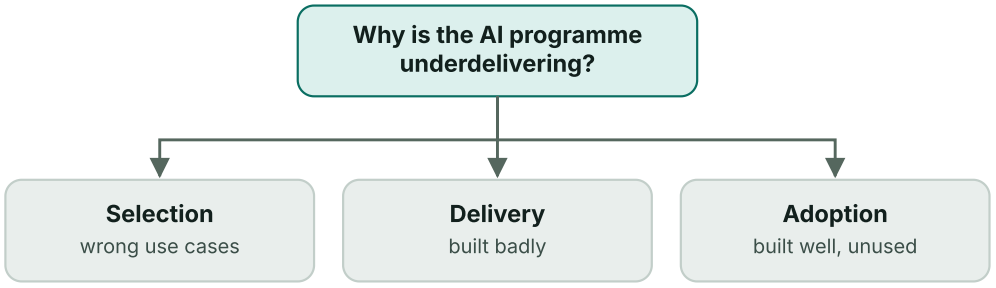
1. Hypothesis-first problem solving

The amateur gathers everything and hopes insight emerges. The consultant states a **testable best guess immediately** ('AI spend tripled: my hypothesis is untiered routing plus no caching'), structures the analysis to confirm or kill it fast, and **updates out loud** when data disagrees. The loop (stated, tested, revised) is the performance, whether or not the first guess was right.

2. Issue trees and MECE

- Decompose the problem into branches that are **mutually exclusive** (no overlap) and **collectively exhaustive** (nothing missing)
- Example: 'Why is the AI programme underdelivering?' → **selection** (wrong use cases) / **delivery** (built badly) / **adoption** (built well, unused)
- Pressure test: can you say what each branch excludes?

✎ DIAGRAM

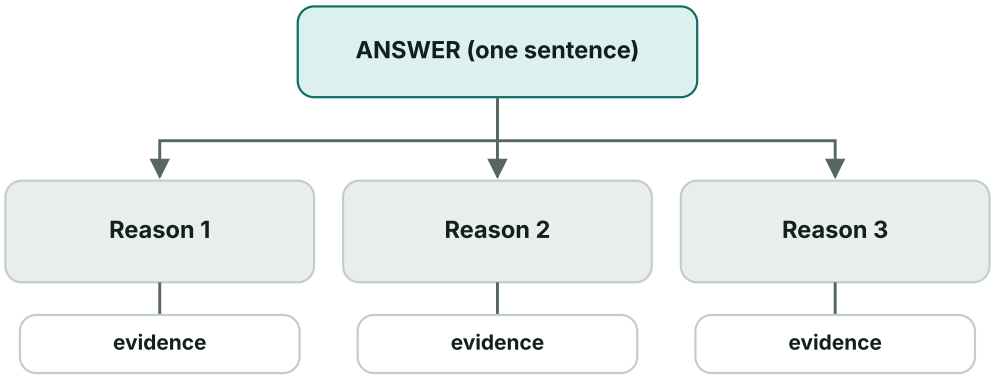


A MECE issue tree: no overlap, nothing missing.

3. The Pyramid Principle

- **Answer first**, then grouped reasons, then evidence
- 'We recommend pausing the rollout: adoption is stalled at 9%, the root cause is workflow mismatch, and a four-week redesign recovers the business case.' Then the detail
- In an interview: one-sentence answer, then 'for three reasons', then give them

✎ DIAGRAM



Executives read the first line and decide whether to read the second.

4. Discovery: value × feasibility × risk

- **Value**: in the client's own arithmetic (volume × time × loaded cost), with a named owner
- **Feasibility**: data readiness, integration surface, pattern maturity, change surface
- **Risk**: cost of an error, regulatory class, reversibility
- The portfolio: 2 to 3 **quick wins** (fund credibility), 1 to 2 **strategic builds**, a **watch list**, and a **kill list** with reasons

✎ DIAGRAM



Start top-right; the first project's job is to earn trust.

5. The three-altitude drill

Explain the same reality to three audiences without condescending:

- **CFO**: money and control ('you're billed by the word, and words written back cost ~5× words read, so tightening verbosity halves the bill')
- **Data lead**: drivers and levers ('output tokens dominate; cut output, cache the prefix, route easy queries')

- **Engineer**: mechanism and configuration ('decode reads the full KV cache per token; set max_tokens, move stable content above the cache breakpoint')

SAY IT

"My hypothesis is X. To test it I'd look at A and B; if A comes back wrong, I'd switch to Y."